Scoring method

We marked each checklist item as: fully met (Yes, green), partially met (Partial, yellow), not met (No, pink/red) or not applicable (N/A, no color).

An item counts as Partial when the information exists but is incomplete, contradicts itself (for example the two different α_max values in the paper), or can only be found in the released code and not in the paper.

N/A items are removed from the count: items that cannot apply to this study, like theoretical proofs, collection of new data or cross-validation.

Each ‚Yes' counts as 1 point, each ‚Partial' as 0.5 and each ‚No' as 0. The score is the total points divided by the number of all (applicable) items.

We report: 67% for NeurIPS (12 of 18) anf 64% for Delphi (16 of 25 items).

The Machine Learning Reproducibility Checklist (v2.0, Apr.7 2020)

For all **models** and **algorithms** presented, check if you include:

- ❑ A clear description of the mathematical setting, algorithm, and/or model.
- ❑ A clear explanation of any assumptions. SAM uncertainity is surrogate, method works on 2D
- ❑ An analysis of the complexity (time, space, sample size) of any algorithm.

there are parameter counts, but no tmings

For any **theoretical claim**, check if you include:

- ❑ A clear statement of the claim.
- ❑ A complete proof of the claim.

For all **datasets** used, check if you include:

- ❑ The relevant statistics, such as number of examples.
- ❑ The details of train / validation / test splits.

preprocessing is applied,
not really explained
and the 256 resolutions missing

- ❑ An explanation of any data that were excluded, and all pre-processing step.
- ❑ A link to a downloadable version of the dataset or simulation environment.
- ❑ For new data collected, a complete description of the data collection process, such as instructions to annotators and methods for quality control.

For all shared **code** related to this work, check if you include:

- ❑ Specification of dependencies.
- ❑ Training code.
- ❑ Evaluation code. inference.py crashed due to the resize problem
- ❑ (Pre-)trained model(s).
- ❑ README file includes table of results accompanied by precise command to run to produce those results.

README has results for the lowest budget,
and generic commands

For all reported **experimental results**, check if you include:

- ❑ The range of hyper-parameters considered, method to select the best hyper-parameter configuration, and specification of all hyper-parameters used to generate results.
- ❑ The exact number of training and evaluation runs.

ranges are given, but a
selection method is not clear

- ❑ A clear definition of the specific measure or statistics used to report results.
- ❑ A description of results with central tendency (e.g. mean) & variation (e.g. error bars).
- ❑ The average runtime for each result, or estimated energy cost.
- ❑ A description of the computing infrastructure used.

Reproduced from: *www.cs.mcgill.ca/~jpineau/ReproducibilityChecklist-v2.0.pdf*

Table 2 Finalized checklist for reproducible reporting of deep learning studies with 26 items

Item	Check	Item	Check
1. Dataset name		2. Dataset owner (how to access the dataset)	
3. Time coverage of data		4. Dataset size (number of patients, studies, and exams)	
5. Inclusion/exclusion criteria in the code, not in paper		6. Data pre-processing steps (data normalization/standardization techniques) as described above	
7. Software/packages used for pre-processing (including version)		8. Transformations (augmentation techniques) (qualitative) (e.g., rotation, affine transformation, Gaussian noise)	
9. Transformations (augmentation techniques) (quantitative) (e.g., degree of rotation, amount of translation, the mean and standard deviation of Gaussian noise) code only		10. Grouping criteria for data splitting (e.g., whether samples are grouped by patient or study)	
11. Stratification criteria		12. Details of cross-validation, nested cross-validation, or similar techniques	
13. The source code for the architecture(s) used [alternatively: denoting the package they used and their version]		14. Hyperparameters for dynamic components of the model's architecture (batch norm layers, dropout probability, etc.) [if used something other than the default value]	
15. Usage of pre-trained weights, source of pre-training weights (if weights are pre-trained externally) and pre-training protocol (if weights are pre-trained internally)		16. Training procedure: loss function(s) and their hyperparameter contradiction on alpha_max	
17. Optimizer(s), learning rate, EMA, and their hyperparameter		18. Learning rate scheduler(s) and their hyperparameter (how learning rate changes, if applicable) code only	
19. Software/packages used for training and statistical analysis (including version) no PyTorch version		20. Number of epochs	
21. Inference procedure: data pre-processing steps (if different from training) crash on the 256 resize		22. Data post-processing steps (if any) some steps only in code	
23. Whether the difference between sets (or folds) has been investigated		24. Metrics, statistical tests are done, evaluation results, and reporting CI for each metric	
25. Definition of ground truth/gold standard only UTAH has expert-annotated labels, no info on other sets		26. Performance of subgroup analysis (e.g., age, sex, disease grade)	